

# RWKV for the Web

Web-RWKV 推理引擎架构解析

web-rwkv Public

Unpin Unwatch 6 Fork 22 Star 282

main 1 Branch 146 Tags

Go to file

Add file

Code

About

Implementation of the RWKV language model in pure WebGPU/Rust.

- Readme
- View license
- Activity
- 282 stars
- 6 watching
- 22 forks

Releases 140

v0.10.10 Latest 2 weeks ago

+ 139 releases


Packages

No packages published  
Publish your first package

Contributors 5



Languages



cryscan	Simplify resource binding APIs. ✓	202c5dc · last week	657 Commits
.cargo	Allow model to be created from bytes.	2 years ago	
.github/workflows	Bump version to v0.10.8	3 weeks ago	
assets	Add RWKV-Othello example (rwkv v7)	last month	
crates/web-rwkv-derive	Make clippy happy.	3 months ago	
examples	Simplify resource binding APIs.	last week	
screenshots	Add screenshots.	2 years ago	
src	Simplify resource binding APIs.	last week	
.gitattributes	Subgroup (#27)	10 months ago	
.gitignore	Update .gitignore	2 months ago	
Cargo.toml	Separate tokio feature completely so that no-default-featur...	2 weeks ago	
LICENSE	Add license.	2 years ago	
README.md	Bump version to v0.10.8	3 weeks ago	

README License

# Web-RWKV

# 为什么是『Web』

一份代码，多端通用

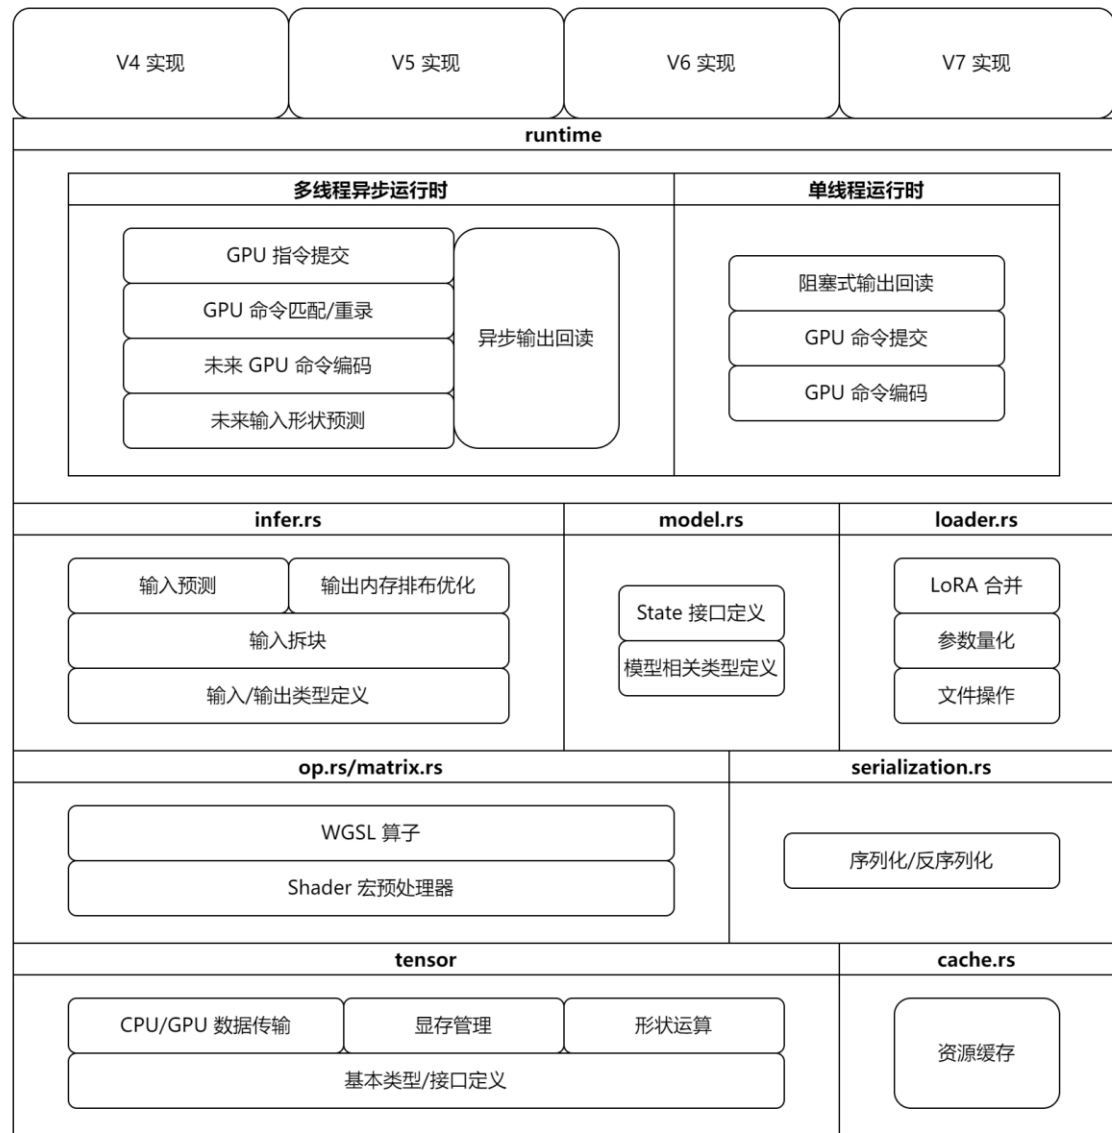
- PC
- Linux
- MacOS
- Android
- iOS
- 浏览器 (WebGPU)



# 特点

- 
- 高度集成：无 Python、CUDA 依赖
  - 广泛兼容：支持 NVIDIA/AMD/Intel GPUs
  - 并行推理：同时批量处理不定长的输入
  - 相对快速：全量 V7 3B 输入 560 t/s 生成 80 t/s @ 3945WX + 3080
  - 静态量化：支持 INT8/NF4/SF4、支持量化参数序列化/反序列化
  - 微调外挂：支持 LoRA、支持 State 热切换
  - 自由灵活：注入 Hook 读写任意推理阶段的数据，修改模型计算

# 架构



# 迭代



## Version 0.1

- 成功运行 RWKV 4 “Dove”




## Version 0.2

- 代码模块化
- 实现 Tensor API
- 实现并行推理
- Crates.io 发布



## Version 0.3

- RWKV 5 “Eagle”
- LoRA 加载



## Version 0.4

- RWKV 6 “Finch”



# 迭代



## Version 0.5

- Shader 宏预处理器
- 同时支持 Fp32/Fp16 读写



## Version 0.8

- 异步运行时
- 多线程 GPU 命令编码
- GPU 命令预测编码
- 使用 Subgroups 操作
- 生成提速 50%-100%



## Version 0.9

- RWKV 7 "Goose"
- 支持在大量设备的浏览器上运行

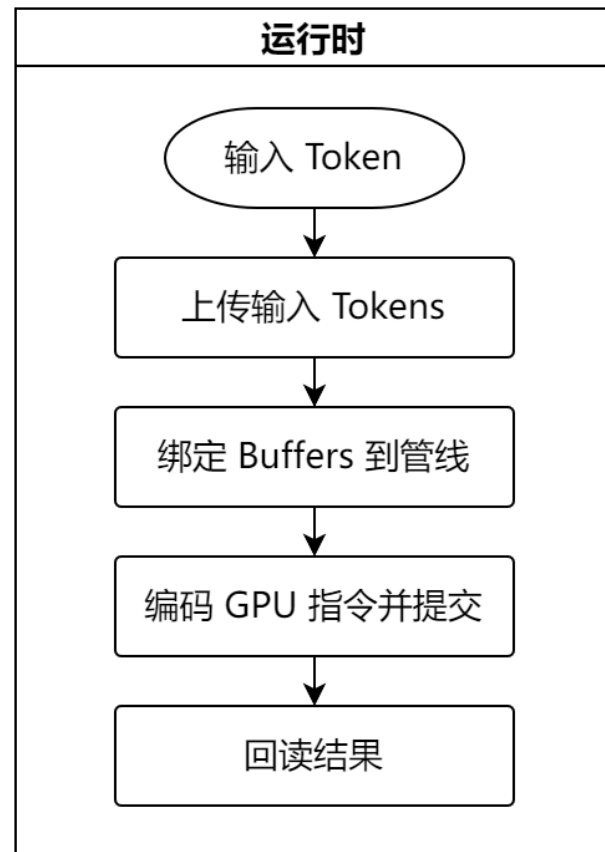
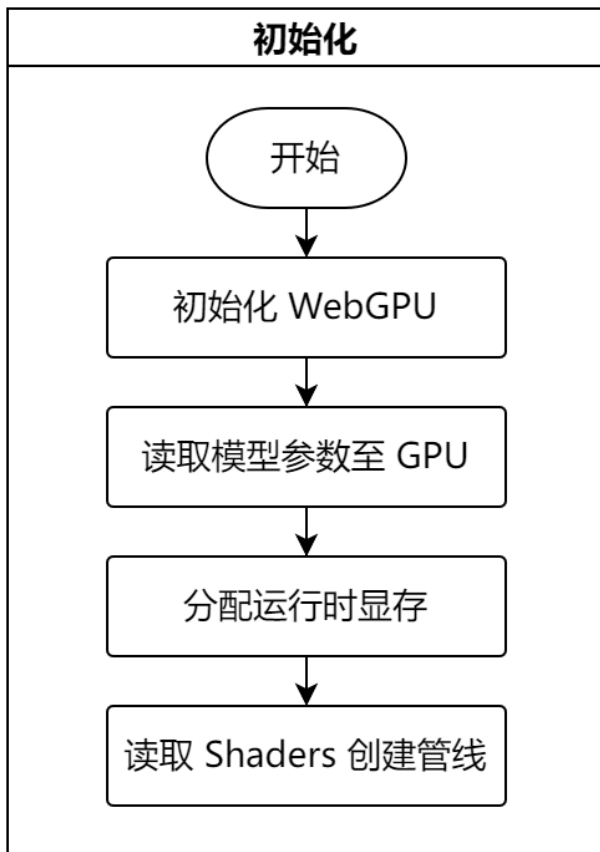


## Version 0.10

- Bind Group 缓存, 生成提速 30%-50%
- 简化 Bind Group 相关 API, 实现模型更方便



# Mini Web-RWKV





## 1. 初始化 WebGPU

```
pub struct Environment {  
    pub instance: Instance,  
    pub adapter: Adapter,  
    pub device: Device,  
    pub queue: Queue,  
}  
  
impl Environment {  
    pub async fn create() -> Result<Self> {  
        let instance = wgpu::Instance::new(InstanceDescriptor::default());  
        let adapter = instance  
            .request_adapter(&RequestAdapterOptions::default())  
            .await?;  
        let (device, queue) = adapter  
            .request_device(&DeviceDescriptor::default(), None)  
            .await?;  
    }  
}
```

```
pub struct Model {
    pub env: Arc<Environment>,
    pub info: ModelInfo,
    pub tensor: Arc<ModelTensor>,
    pub pipeline: Arc<ModelPipeline>,
}

pub struct ModelInfo {
    pub num_layers: usize,
    pub num_emb: usize,
    pub num_vocab: usize,
}

pub struct ModelTensor {
    pub dim: Buffer,
    pub embed: Embed,
    pub head: Head,
    pub layers: Vec<Layer>,
}
```

```
pub struct LayerNorm {
    pub w: Buffer,
    pub b: Buffer,
}

pub struct Embed {
    pub layer_norm: LayerNorm,
    pub w: Vec<f16>,
}

pub struct Head {
    pub layer_norm: LayerNorm,
    pub dims: Buffer,
    pub w: Buffer,
}
```

```
pub struct Att {
    pub time_decay: Buffer,
    pub time_first: Buffer,
    pub dims: Buffer,
    pub time_mix_k: Buffer,
    pub time_mix_v: Buffer,
    pub time_mix_r: Buffer,
    pub w_k: Buffer,
    pub w_v: Buffer,
    pub w_r: Buffer,
    pub w_o: Buffer,
}

pub struct Ffn {
    pub time_mix_k: Buffer,
    pub time_mix_r: Buffer,
    pub dims_k: Buffer,
    pub dims_v: Buffer,
    pub dims_r: Buffer,
    pub w_k: Buffer,
    pub w_v: Buffer,
    pub w_r: Buffer,
}
```

## 2. 声明模型结构

```
pub struct ModelBuffer {  
    pub tokens: Vec<u16>,  
    pub num_tokens: Buffer,  
  
    pub emb_x: Buffer,  
    pub emb_o: Buffer,  
  
    pub att_x: Buffer,  
    pub att_kx: Buffer,  
    pub att_vx: Buffer,  
    pub att_rx: Buffer,  
    pub att_k: Buffer,  
    pub att_v: Buffer,  
    pub att_r: Buffer,  
    pub att_w: Buffer,  
    pub att_o: Buffer,  
  
    pub ffn_x: Buffer,  
    pub ffn_kx: Buffer,  
    pub ffn_vx: Buffer,  
    pub ffn_rx: Buffer,  
    pub ffn_k: Buffer,  
    pub ffn_v: Buffer,  
    pub ffn_r: Buffer,  
    pub ffn_o: Buffer,  
  
    pub head_x: Buffer,  
    pub head_r: Buffer,  
    pub head_o: Buffer,  
  
    pub map: Buffer,  
}  
  
pub struct ModelState(pub Vec<LayerState>);  
  
pub struct LayerState {  
    pub att: Buffer,  
    pub ffn: Buffer,  
}
```

### 3. 声明运行时需要的 Buffer

#### 4. 从文件中读取模型参数，分配参数显存并上传

```
impl Model {
    fn from_bytes(data: &[u8], env: Arc<Environment>) -> Result<Self> {
        let device = &env.device;
        let model = SafeTensors::deserialize(data)?;
        let load_tensor_f32 = |name: String| -> Result<Buffer> {
            let tensor = model.tensor(&name)?.data();
            let tensor: Vec<_> = pod_collect_to_vec::<_, f16>(tensor)
                .into_iter()
                .map(f16::to_f32)
                .collect();
            let buffer = device.create_buffer_init(&BufferInitDescriptor {
                label: Some(&name),
                contents: cast_slice(&tensor),
                usage: BufferUsages::STORAGE,
            });
            Ok(buffer)
        };
        let load_tensor_f16 = |name: String| -> Result<Buffer> {
            let tensor = model.tensor(&name)?.data();
            let buffer = device.create_buffer_init(&BufferInitDescriptor {
                label: Some(&name),
                contents: cast_slice(tensor),
                usage: BufferUsages::STORAGE,
            });
            Ok(buffer)
        };
        let embed = Embed {
            layer_norm: LayerNorm {
                w: load_tensor_f32("blocks.0.ln0.weight".into())?,
                b: load_tensor_f32("blocks.0.ln0.bias".into())?,
            },
            w: pod_collect_to_vec(model.tensor("emb.weight")?.data()),
        };
        let head = Head {
            layer_norm: LayerNorm {
                w: load_tensor_f32("ln_out.weight".into())?,
                b: load_tensor_f32("ln_out.bias".into())?,
            },
            dims: create_uniform_u32(&[num_emb as u32, num_vocab as u32]),
            w: load_tensor_f16("head.weight".into())?,
        };
    }
}
```

## 5. Buffer 分配显存

```
pub fn create_buffer(&self, tokens: &[u16]) -> ModelBuffer {
    let device = &self.env.device;

    let create_buffer_f32 = |capacity: usize| -> Buffer {
        let data = vec![0.0f32; capacity];
        device.create_buffer_init(&BufferInitDescriptor {
            label: None,
            contents: cast_slice(&data),
            usage: BufferUsages::STORAGE | BufferUsages::COPY_DST | BufferUsages::COPY_SRC,
        })
    };

    let load_buffer_f32 = |data: &[f32]| -> Buffer {
        device.create_buffer_init(&BufferInitDescriptor {
            label: None,
            contents: cast_slice(data),
            usage: BufferUsages::STORAGE | BufferUsages::COPY_DST,
        })
    };

    let create_uniform_u32 = |values: &[u32]| -> Buffer {
        device.create_buffer_init(&BufferInitDescriptor {
            label: None,
            contents: cast_slice(values),
            usage: BufferUsages::UNIFORM | BufferUsages::COPY_DST,
        })
    };
}
```

```
let num_tokens = tokens.len();
let num_emb = self.info.num_emb;
let num_vocab = self.info.num_vocab;
let capacity = num_tokens * num_emb;

let input = self.embedding(tokens);

let map = device.create_buffer(&BufferDescriptor {
    label: None,
    size: 4 * num_vocab as u64,
    usage: BufferUsages::MAP_READ | BufferUsages::COPY_DST,
    mapped_at_creation: false,
});

ModelBuffer {
    tokens: tokens.to_vec(),
    num_tokens: create_uniform_u32(&[num_tokens as u32]),
    emb_x: load_buffer_f32(&input),
    emb_o: create_buffer_f32(capacity),
    att_x: create_buffer_f32(capacity),
    att_kx: create_buffer_f32(capacity),
    att_vx: create_buffer_f32(capacity),
    att_rx: create_buffer_f32(capacity),
    att_k: create_buffer_f32(capacity),
    att_v: create_buffer_f32(capacity),
    att_r: create_buffer_f32(capacity),
    att_w: create_buffer_f32(capacity),
    att_o: create_buffer_f32(capacity),
    ffn_x: create_buffer_f32(capacity),
    ffn_kx: create_buffer_f32(capacity),
    ffn_vx: create_buffer_f32(4 * capacity),
    ffn_rx: create_buffer_f32(capacity),
    ffn_k: create_buffer_f32(4 * capacity),
    ffn_v: create_buffer_f32(capacity),
    ffn_r: create_buffer_f32(capacity),
    ffn_o: create_buffer_f32(capacity),
    head_x: create_buffer_f32(num_emb),
    head_r: create_buffer_f32(num_emb),
    head_o: create_buffer_f32(num_vocab),
    map,
}
```

## 5. State 分配显存

```
pub fn create_state(&self) -> ModelState {
    let device = &self.env.device;

    let ModelInfo {
        num_layers,
        num_emb,
        ..
    } = self.info;

    let create_buffer_f32 = |data: &[f32]| -> Buffer {
        device.create_buffer_init(&BufferInitDescriptor {
            label: None,
            contents: cast_slice(data),
            usage: BufferUsages::STORAGE | BufferUsages::COPY_DST | BufferUsages::COPY_SRC,
        })
    };

    let mut layers = vec![];
    for _ in 0..num_layers {
        let mut att = vec![0.0f32; 4 * num_emb];
        att[3 * num_emb..4 * num_emb]
            .iter_mut()
            .for_each(|x| *x = -1.0e30);

        let ffn = vec![0.0f32; num_emb];

        let layer = LayerState {
            att: create_buffer_f32(&att),
            ffn: create_buffer_f32(&ffn),
        };
        layers.push(layer);
    }

    ModelState(layers)
}
```

```

@group(0) @binding(0) var<uniform> dims: vec2<u32>;           // [C, R]
@group(0) @binding(1) var<storage, read> matrix: array<vec2<u32>>; // (R, C)
@group(0) @binding(2) var<storage, read> input: array<vec4<f32>>; // (T, C)
@group(0) @binding(3) var<storage, read_write> output: array<vec4<f32>>; // (T, R)

```

```

const BLOCK_SIZE: u32 = 256u;

```

```

var<workgroup> local_sum: array<vec4<f32>, BLOCK_SIZE>;

```

```

fn reduce_step_barrier(index: u32, stride: u32) {
    if index < stride {
        local_sum[index] += local_sum[index + stride];
    }
    workgroupBarrier();
}

```

```

@compute @workgroup_size(256, 1, 1)
fn matmul(@builtin(global_invocation_id) invocation_id: vec3<u32>) {
    let index = invocation_id.x;
    let channel = invocation_id.y;           // 1 channel: 4 rows in matrix
    let token = invocation_id.z;
    let stride = dims / 4u;

```

```

    local_sum[index] = vec4<f32>(0.0);
    for (var i = index; i < stride.x; i += BLOCK_SIZE) {
        let ti = token * stride.x + i;
        var ci = channel * 4u * stride.x + i;

```

```

        // read 4 elements from the input
        let x = input[ti];

```

```

        // read 4 rows from the matrix, each with 4 unpacked floats, forming a 4x4 sub-block
        var data: vec2<u32>;
        var m: mat4x4<f32>;

```

```

        data = matrix[ci]; m[0] = vec4<f32>(unpack2x16float(data.x), unpack2x16float(data.y)); ci += stride.x;
        data = matrix[ci]; m[1] = vec4<f32>(unpack2x16float(data.x), unpack2x16float(data.y)); ci += stride.x;
        data = matrix[ci]; m[2] = vec4<f32>(unpack2x16float(data.x), unpack2x16float(data.y)); ci += stride.x;
        data = matrix[ci]; m[3] = vec4<f32>(unpack2x16float(data.x), unpack2x16float(data.y));
        local_sum[index] += transpose(m) * x;

```

```

    }
    workgroupBarrier();

```

```

    reduce_step_barrier(index, 128u);
    reduce_step_barrier(index, 64u);
    reduce_step_barrier(index, 32u);

```

```

    if index < 32u {
        local_sum[index] += local_sum[index + 16u];
        local_sum[index] += local_sum[index + 8u];
        local_sum[index] += local_sum[index + 4u];
        local_sum[index] += local_sum[index + 2u];
        local_sum[index] += local_sum[index + 1u];
    }

```

```

    if index == 0u {
        output[token * stride.y + channel] = local_sum[0];
    }
}

```

## 6. 手写 Compute Shader，实现

- Layer Normalization
- 矩阵×向量
- Time Shift
- WKV Operator
- Channel Mix



## 7. 读取 Shader, 创建 Compute Pipeline

```
https://tool.geekfa.com/

let create_pipeline = |shader: &str, entry_point: &str| -> ComputePipeline {
  let module = &device.create_shader_module(ShaderModuleDescriptor {
    label: None,
    source: ShaderSource::Wgsl(Cow::Borrowed(shader)),
  });
  device.create_compute_pipeline(&ComputePipelineDescriptor {
    label: None,
    layout: None,
    module,
    entry_point,
  })
};

let pipeline = Arc::new(ModelPipeline {
  layer_norm: create_pipeline(include_str!("shaders/layer_norm.wgsl"), "layer_norm"),
  token_shift: create_pipeline(include_str!("shaders/token_shift.wgsl"), "token_shift"),
  matmul: create_pipeline(include_str!("shaders/matmul.wgsl"), "matmul"),
  token_mix: create_pipeline(include_str!("shaders/token_mix.wgsl"), "token_mix"),
  activation: create_pipeline(include_str!("shaders/activation.wgsl"), "activation"),
  channel_mix: create_pipeline(include_str!("shaders/channel_mix.wgsl"), "channel_mix"),
  add: create_pipeline(include_str!("shaders/add.wgsl"), "add"),
});
```

```
https://tool.geekfa.com/

pub struct ModelPipeline {
  pub layer_norm: ComputePipeline,
  pub token_shift: ComputePipeline,
  pub matmul: ComputePipeline,
  pub token_mix: ComputePipeline,
  pub activation: ComputePipeline,
  pub channel_mix: ComputePipeline,
  pub add: ComputePipeline,
}
```



## 8. 当输入了一个 Token: 转化为 Embedding 并上传 GPU

```
pub fn embedding(&self, tokens: &[u16]) -> Vec<f32> {  
    let num_tokens = tokens.len();  
    let num_emb = self.info.num_emb;  
    let capacity = num_tokens * num_emb;  
  
    let mut input = vec![];  
    input.reserve(capacity);  
    for token in tokens {  
        let index = *token as usize;  
        let mut embed: Vec<_> = self.tensor.embed.w[index * num_emb..(index + 1) * num_emb]  
            .iter()  
            .copied()  
            .map(f16::to_f32)  
            .collect();  
        input.append(&mut embed);  
    }  
    input  
}
```

```

fn create_bind_group(&self, buffer: &ModelBuffer, state: &ModelState) -> ModelBindGroup {
    let device = &self.env.device;
    let pipeline = &self.pipeline;

    let [layer_norm_layout, token_shift_layout, matmul_layout, token_mix_layout, activation_layout, channel_mix_layout,
        [
            &pipeline.layer_norm,
            &pipeline.token_shift,
            &pipeline.matmul,
            &pipeline.token_mix,
            &pipeline.activation,
            &pipeline.channel_mix,
            &pipeline.add,
        ]
    ].map(|pipeline| pipeline.get_bind_group_layout(0));

    let embed = {
        let layer_norm = device.create_bind_group(&BindGroupDescriptor {
            label: None,
            layout: &layer_norm_layout,
            entries: &[
                BindGroupEntry {
                    binding: 0,
                    resource: buffer.emb_x.as_entire_binding(),
                },
                BindGroupEntry {
                    binding: 1,
                    resource: buffer.emb_x.as_entire_binding(),
                },
                BindGroupEntry {
                    binding: 2,
                    resource: self.tensor.embed.layer_norm.w.as_entire_binding(),
                },
                BindGroupEntry {
                    binding: 3,
                    resource: self.tensor.embed.layer_norm.b.as_entire_binding(),
                },
                BindGroupEntry {
                    binding: 4,
                    resource: buffer.emb_o.as_entire_binding(),
                },
            ],
        });
        EmbedBindGroup { layer_norm }
    };
};

```

## 9. 绑定 Buffers 到 Pipeline; 创建 Compute Pass; 编码、提交 GPU 命令

```

fn run_internal(&self, buffer: &ModelBuffer, state: &ModelState) {
    let device = &self.env.device;
    let queue = &self.env.queue;

    let bind_group = self.create_bind_group(buffer, state);
    let pipeline = &self.pipeline;

    let ModelInfo {
        num_emb, num_vocab, ..
    } = self.info;

    let num_tokens = buffer.tokens.len() as u32;
    let num_emb_vec4 = num_emb as u32 / 4;
    let num_vocab_vec4 = num_vocab as u32 / 4;
    const BLOCK_SIZE: u32 = 256;

    let mut encoder = device.create_command_encoder(&CommandEncoderDescriptor::default());
    {
        let mut pass = encoder.begin_compute_pass(&ComputePassDescriptor::default());
        pass.set_pipeline(&pipeline.layer_norm);
        pass.set_bind_group(0, &bind_group.embed.layer_norm, &[]);
        pass.dispatch_workgroups(1, num_tokens, 1);
    }
    for layer in &bind_group.layers {
        {
            let mut pass = encoder.begin_compute_pass(&ComputePassDescriptor::default());
            pass.set_pipeline(&pipeline.layer_norm);
            pass.set_bind_group(0, &layer.att_layer_norm, &[]);
            pass.dispatch_workgroups(1, num_tokens, 1);
            // ...
        }
        // ...
        encoder.copy_buffer_to_buffer(&buffer.head_o, 0, &buffer.map, 0, 4 * num_vocab as u64);
        queue.submit(Some(encoder.finish()));
    }
}

```

```
https://tool.geekfa.com/

let (sender, receiver) = async_channel::bounded(1);
let slice = buffer.map.slice(..);
slice.map_async(wgpu::MapMode::Read, move |v| {
    sender.send_blocking(v).unwrap();
});
// 等待 GPU 计算完成, 阻塞!
self.env.device.poll(wgpu::MaintainBase::Wait);
match receiver.recv_blocking() {
    Ok(_) => {
        let data = {
            let data = slice.get_mapped_range();
            cast_slice(&data).to_vec()
        };
        buffer.map.unmap();
        Ok(data)
    }
    Err(err) => Err(err.into()),
}
```

## 10. 回读计算结果

# 恭喜你实现了 Web-RWKV 0.1

- 
- 速度比 Torch 实现快一些……
  - V4 3B 50 t/s @ 3945WX + 3080
  - 但是如果我的机器是 Xeon E5-2673 v3 + 4090
  - 那么速度只有 25 t/s (……)
  - 我们一步一步来优化它……

# 长输入优化及并行推理


- 
- 长输入使用 GEMM 而非 GEMV，节省大量带宽，提速 8-10 倍
  - 为了效率，GEMM Kernel 期望输入序列长是 32 的倍数
  - 过长的输入可能会使 GPU 长时间繁忙从而触发 TDR
  - 故我们需要将长输入拆分成合理大小的块，一次处理一块
  - 那并行推理也可以复用这个加速机制吗？

# 从推理函数的类型中我们可以看到……




<https://tool.geekfa.com/>

```
/// Perform (partial) inference and return the remaining input and (perhaps partial) output.  
/// The amount of input processed during one call is bound by the input chunk size.  
pub async fn infer(&self, input: I) -> Result<(I, O)> {  
    let (sender, receiver) = flume::bounded(1);  
    let submission = Submission { input, sender };  
    let _ = self.0.send(submission).await;  
    receiver.recv_async().await?  
}
```



## WKV Kernel




Input

1 Chunk

## Output

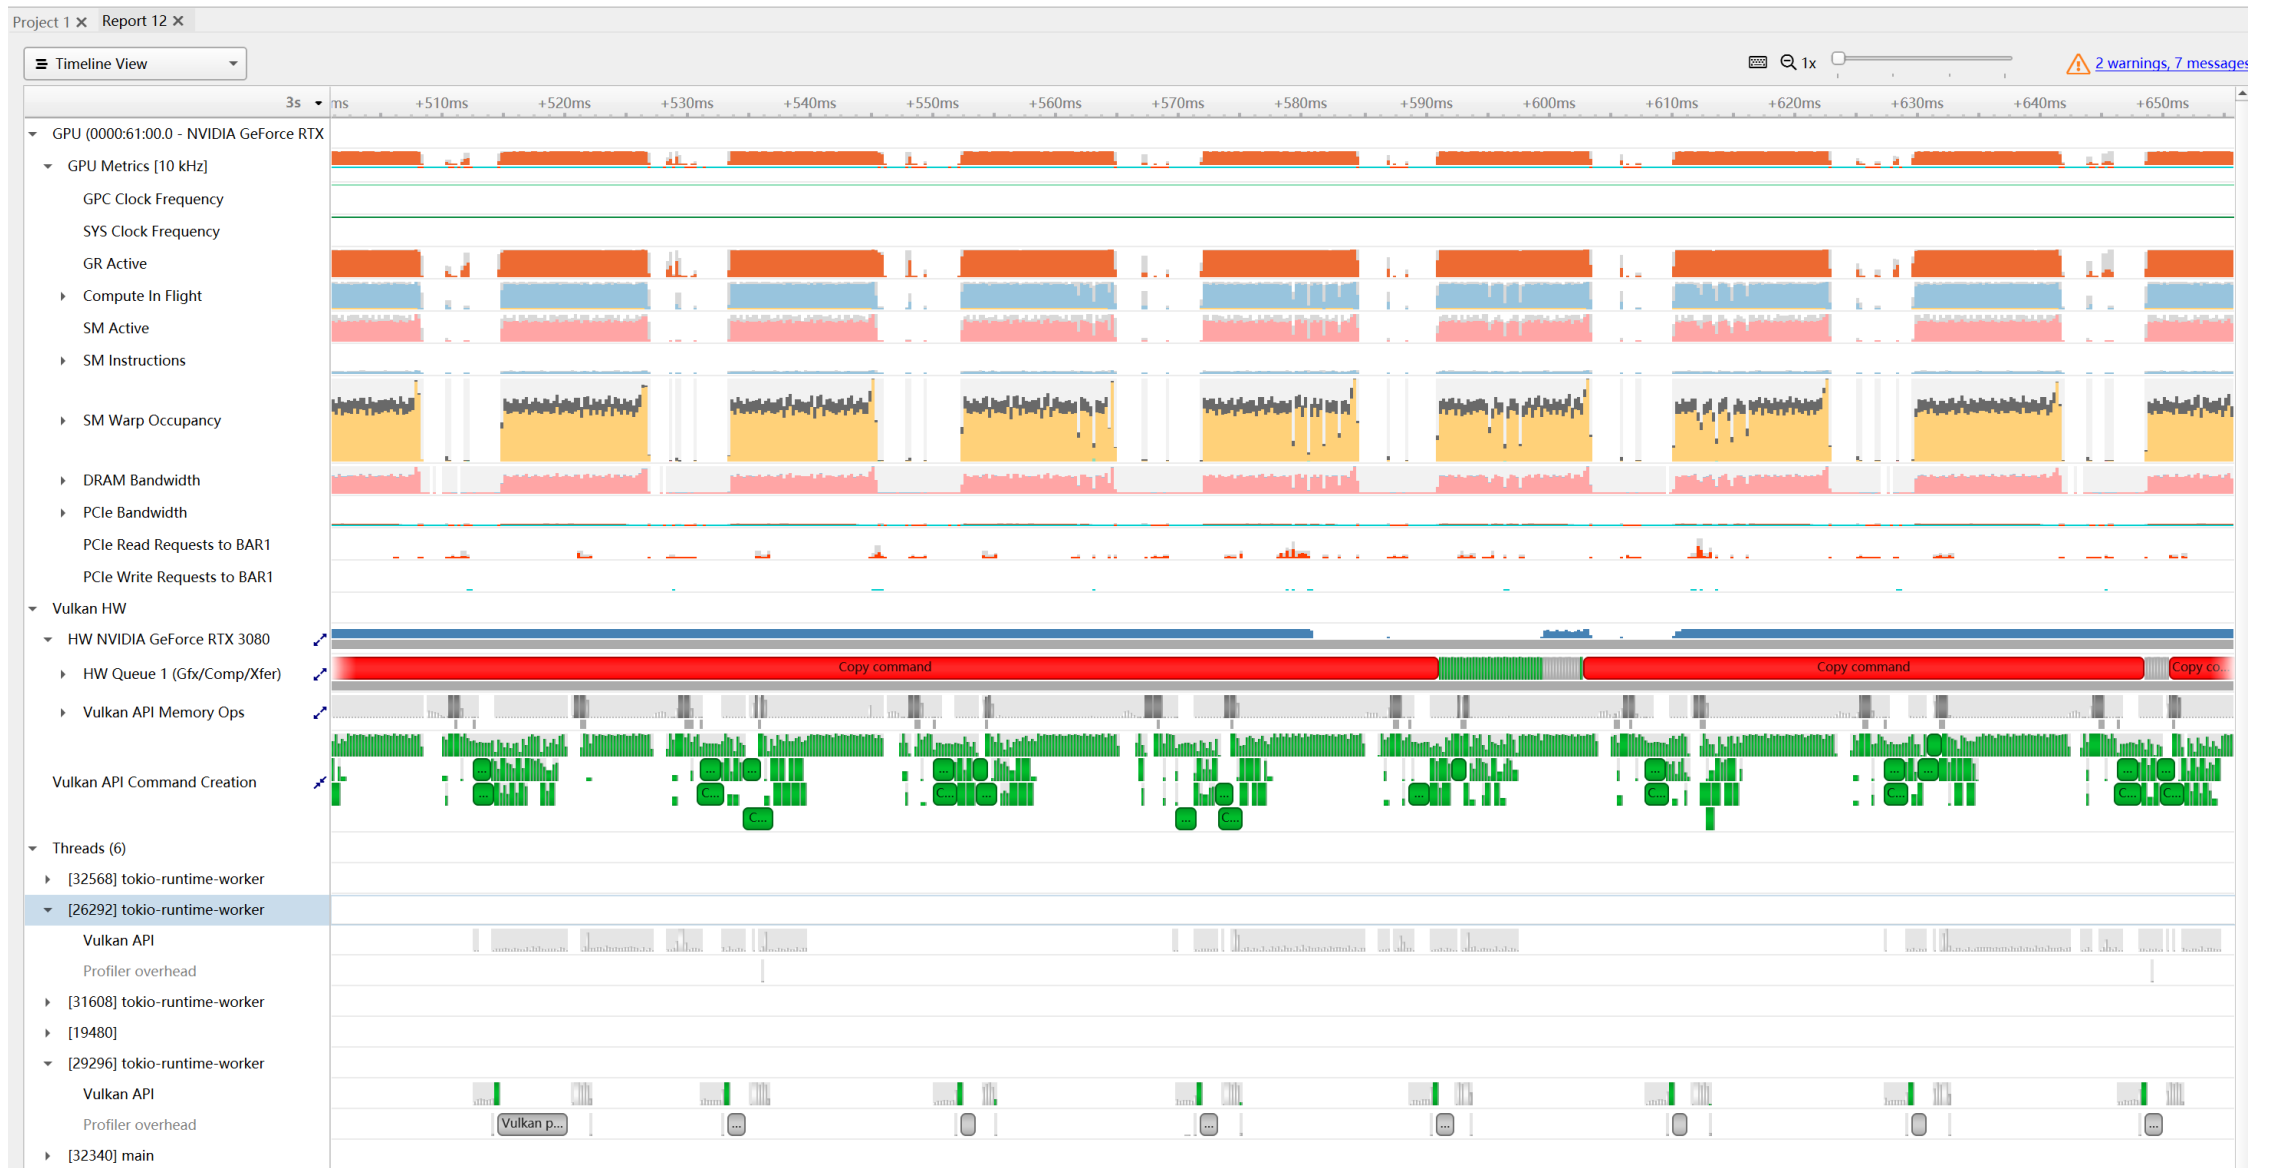
## Infer Calls



An illustration of a modern office environment. In the foreground, a person is seen from behind, sitting at a desk and working on a laptop. The desk is cluttered with papers, a glass of water, and other office supplies. In the background, several other people are standing and talking, some holding devices. Large windows on the left side of the room let in bright light, and abstract art pieces are hanging on the wall.

**输入快了不少!  
可是生成速度  
还是慢……做个  
性能测试吧?**

推理引擎内置 Tracy 性能测试接口,  
只要开启相应的编译特性即可。



AMD Ryzen Threadripper PRO 3945WX + 3080

# GPU 在等什么?

```
let att_matmul_k = device.create_bind_group(&BindGroupDescriptor {
  label: None,
  layout: &matmul_layout,
  entries: &[
    BindGroupEntry {
      binding: 0,
      resource: layer.att.dims.as_entire_binding(),
    },
    BindGroupEntry {
      binding: 1,
      resource: layer.att.w_k.as_entire_binding(),
    },
    BindGroupEntry {
      binding: 2,
      resource: buffer.att_kx.as_entire_binding(),
    },
    BindGroupEntry {
      binding: 3,
      resource: buffer.att_k.as_entire_binding(),
    },
  ],
});
```

## 创建绑定组

```
let mut pass = encoder.begin_compute_pass(&ComputePassDescriptor {
  label: None,
});
pass.set_pipeline(&pipeline.layer_norm);
pass.set_bind_group(0, &bind_group.head.layer_norm, &[]);
pass.dispatch_workgroups(1, 1, 1);

pass.set_pipeline(&pipeline.matmul);
pass.set_bind_group(0, &bind_group.head.matmul, &[]);
pass.dispatch_workgroups(1, num_vocab_vec4, 1);


encoder.copy_buffer_to_buffer(&buffer.head_o, 0, &buffer.map,
```

## 编码命令




# 绑定组创建优化

多层次资源缓存复用



# 绑定组创建优化








命令编码优化

# 输入预测

- 
- 在生成上一输出前，无法得知下一输入？
  - 至少我们可以预测形状（假设：推理过程不被人为中断）
  - 在 GPU 计算第  $N$  个 Token 时，CPU 并行编码第  $N + 1$ 、 $N + 2$ ……
  - 如果预测错误，则停止现有编码任务，冲刷流水线，重新编码
  - 但是，CPU 要如何生产才能达到供需平衡？

- 生产者 (CPU)  $\leftrightarrow$  消费者 (GPU)
- 平衡方程
- 动态调整 CPU 编码任务数
- $2 \rightarrow 1 \rightarrow 0$  循环

Xeon E5, 相比同步编码, 异步并行预测编码提速 100%





WKV 5/6/7

				T0				T1				T2				T3				A/B/R/W/K					
				↓				↓				↓				↓									
S					α0	β0	γ0	δ0	ε0	ζ0	η0	θ0	ι0	κ0	λ0	μ0	ν0	ξ0	ο0	π0	Loop/Reduction Direction ↓				0
					α1	β1	γ1	δ1	ε1	ζ1	η1	θ1	ι1	κ1	λ1	μ1	ν1	ξ1	ο1	π1					1
					α2	β2	γ2	δ2	ε2	ζ2	η2	θ2	ι2	κ2	λ2	μ2	ν2	ξ2	ο2	π2					2
					α3	β3	γ3	δ3	ε3	ζ3	η3	θ3	ι3	κ3	λ3	μ3	ν3	ξ3	ο3	π3					3
					α4	β4	γ4	δ4	ε4	ζ4	η4	θ4	ι4	κ4	λ4	μ4	ν4	ξ4	ο4	π4					4
					α5	β5	γ5	δ5	ε5	ζ5	η5	θ5	ι5	κ5	λ5	μ5	ν5	ξ5	ο5	π5					5
					α6	β6	γ6	δ6	ε6	ζ6	η6	θ6	ι6	κ6	λ6	μ6	ν6	ξ6	ο6	π6					6
					α7	β7	γ7	δ7	ε7	ζ7	η7	θ7	ι7	κ7	λ7	μ7	ν7	ξ7	ο7	π7					7
					α8	β8	γ8	δ8	ε8	ζ8	η8	θ8	ι8	κ8	λ8	μ8	ν8	ξ8	ο8	π8					8
					α9	β9	γ9	δ9	ε9	ζ9	η9	θ9	ι9	κ9	λ9	μ9	ν9	ξ9	ο9	π9					9
					αA	βA	γA	δA	εA	ζA	ηA	θA	ιA	κA	λA	μA	νA	ξA	οA	πA					A
					αB	βB	γB	δB	εB	ζB	ηB	θB	ιB	κB	λB	μB	νB	ξB	οB	πB					B
					αC	βC	γC	δC	εC	ζC	ηC	θC	ιC	κC	λC	μC	νC	ξC	οC	πC					C
					αD	βD	γD	δD	εD	ζD	ηD	θD	ιD	κD	λD	μD	νD	ξD	οD	πD					D
					αE	βE	γE	δE	εE	ζE	ηE	θE	ιE	κE	λE	μE	νE	ξE	οE	πE					E
					αF	βF	γF	δF	εF	ζF	ηF	θF	ιF	κF	λF	μF	νF	ξF	οF	πF					F
				Memory/Thread Layout →																					
SA/V				α	β	γ	δ	ε	ζ	η	θ	ι	κ	λ	μ	ν	ξ	ο	π						

# 有关平台兼容性的一些坑……

- 
- 不要把一些 N 卡经验带进 Shader（该同步还得同步）
  - A 卡发现 NF4 的 Quantiles 每次从共享内存取一个 f32 更快，欣然提交；一小时后发现 I 卡只可取 `vec4<f32>` 对齐，否则乱码……
  - 苹果芯片的 `tanh` 内置函数输出超过 42 返回 NaN（最诨诨）
  - ……